

# 论文研讨 - LIGHTYEAR: Using Modularity to Scale BGP Control Plane Verification

Alan Tang <sup>1</sup>, Ryan Beckett <sup>2</sup>, Steven Benaloh <sup>2</sup>, Karthick Jayaraman <sup>2</sup>, Tejas Patil <sup>2</sup>, Todd Millstein <sup>1</sup>,  
George Varghese <sup>1</sup>,  
<sup>1</sup> UCLA, <sup>2</sup> Microsoft

# Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

$$\forall r, s \in \mathbb{Q}, r = s \vee \exists t \in \mathbb{Q}, r < t < s$$

# Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions)  $\Rightarrow$  Postconditions (Assertions)

# Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions)  $\Rightarrow$  Postconditions (Assertions)

Two way of reasoning about their correctness:

- Formal deduction
-

# Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions)  $\Rightarrow$  Postconditions (Assertions)

Two way of reasoning about their correctness:

- Formal deduction
- **Model checking**

# Background

A lot of desirable properties are **global**:

- When **all nodes** works, no incorrect routes are advertised.
- When **all edges** works, **all nodes** acknowledge correct routes.

# Background

A lot of desirable properties are **global**:

- When **all nodes** works, no incorrect routes are advertised.
- When **all edges** works, **all nodes** acknowledge correct routes.

Checking global properties **globally** is HARD.

- Symbolically: Amount of states grows exponentially, amount of checks grows (at least) quadratically.
- Runtime: Synchronization is costly / hard to do right.

It's much easier to do **local checks**.



It's much easier to do **local checks**.  
**LIGHTYEAR**

It's much easier to do **local checks**.

## **LIGHTYEAR**

Given:

- Configuration (Topology and policies)
- Local invariants
- Desirable global properties

Generates:

- Local checks

Paper proves that local constraints imply global properties


# Scope of LIGHTYEAR

Internal WAN (with BGP support).

# Scope of LIGHTYEAR

Internal WAN (with BGP support).

**Cloud provider:** Vultr, Azure, ...



# The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).

# The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).

E.g. At inner node  $N$ :

$$\text{addr}(r) \in \text{dead:beef::}/64$$

# The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).


E.g. At inner node  $N$ :

$$\text{addr}(r) \in \text{dead:beef}::/64$$

Generated functions: Import, Export, Originates


- Import:  $\text{Edge} \times \text{Route} \rightarrow \text{Route} \cup \{ \text{Reject} \}$

# An example...






# An example...



## Safety:

No routes from ISP1  
reaches ISP2

# An example...



## Safety:

No routes from ISP1 reaches ISP2

## Liveness

Routes from customer reaches ISP2

# Limits of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

# Limits of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

LIGHTYEAR does not use temporal information. We only care about:

- Safety: No bogus routes
- Liveness: Valid routes will be broadcasted

$$\Diamond\Diamond A \rightarrow \Diamond A$$

# Limits of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

LIGHTYEAR does not use temporal information. We only care about:

- Safety: No bogus routes
- Liveness: Valid routes will be broadcasted

$$\Diamond\Diamond A \rightarrow \Diamond A$$

**Topology and policies are fixed**, and for liveness checks, the path needs to be provided.

**Proof? How?**

**Proof? How?**

**SMT**

**Proof? How?**

# SMT

For all location  $l$ , generates:

$$\left(\bigwedge \text{inv}(l)\right) \wedge (r' = \text{Import}(l, r)) \rightarrow \text{inv}(l)[r/r']$$



# Real world deployments

Hundreds of routers, tens of thousands of peering sessions.

Implemented in C#, using Zen as SMT library.

# Real world deployments

Hundreds of routers, tens of thousands of peering sessions.

Implemented in C#, using Zen as SMT library.

11 configuration bugs. Each run takes at most 15 mins.

# Conclusion

LIGHTYEAR is a **modular** way to verify global properties of BGP control plane.

- Guaranteed correctness for safety and liveness.
- Constraints are simple to write and understand.
- Implemented in C# which can be directly integrated into existing systems.

**Thank you!**

