

论文研讨 - LIGHTYEAR: Using Modularity to Scale BGP Control Plane Verification

Alan Tang ¹, Ryan Beckett ², Steven Benaloh ², Karthick Jayaraman ², Tejas Patil ², Todd Millstein ¹,
George Varghese ¹,
¹ UCLA, ² Microsoft

Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

$$\forall r, s \in \mathbb{Q}, r = s \vee \exists t \in \mathbb{Q}, r < t < s$$

Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions) \Rightarrow Postconditions (Assertions)

Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions) \Rightarrow Postconditions (Assertions)

Two way of reasoning about their correctness:

- Formal deduction
-

Pre-Background...

Properties, represented by **propositions and predicates**, are syntactical objects.

Precondition (Assumptions) \Rightarrow Postconditions (Assertions)

Two way of reasoning about their correctness:

- Formal deduction
- **Model checking**

Background

A lot of desirable properties are **global**:

- When **all nodes** works, no incorrect routes are advertised.
- When **all edges** works, **all nodes** acknowledge correct routes.

Background

A lot of desirable properties are **global**:

- When **all nodes** works, no incorrect routes are advertised.
- When **all edges** works, **all nodes** acknowledge correct routes.

Verifying global properties **globally** is HARD.

- Symbolically: Amount of states grows exponentially, amount of checks grows (at least) quadratically.
- Runtime: Synchronization is costly / hard to do right.

It's much easier to do **local checks**.

It's much easier to do **local checks**.
LIGHTYEAR

It's much easier to do **local checks**.

LIGHTYEAR

Given:

- Configuration (Topology and policies)
- Local invariants
- Desirable global properties

Generates:

- Local checks
- Proof that local constraints imply global properties

The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).

The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).

E.g. At node N :

$$\text{addr}(r) \in \text{dead:beef::}/64$$

The language of LIGHTYEAR

Basically the quantifier-free fragment of FOL, specifying behaviors at **locations** (nodes or edges).


E.g. At node N :

$$\text{addr}(r) \in \text{dead:beef::}/64$$


Generated functions: Import, Export, Originates

- Import: $\text{Edge} \times \text{Route} \rightarrow \text{Route} \cup \{ \text{Reject} \}$

An example...




An example...



Safety:

No routes from ISP1
reaches ISP2

An example...




Safety:

No routes from ISP1 reaches ISP2

Liveness

Routes from customer reaches ISP2

An example...



Property: At $R2 \rightarrow \text{ISP2}$

$\neg \text{FromISP1}(r)$

Constraints: At all internal edges:

$\text{FromISP1}(r)$

$\rightarrow 100 : 1 \in \text{Comm}(r)$

Scope of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

Scope of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

LIGHTYEAR does not use temporal information. We only care about:

- Safety: No bogus routes
- Liveness: Valid routes will be broadcasted

$$\Diamond\Diamond A \rightarrow \Diamond A$$

Scope of LIGHTYEAR

Traditionally, predicates are **indexed by time**, or reasoning is done in **temporal logic**.

LIGHTYEAR does not use temporal information. We only care about:

- Safety: No bogus routes
- Liveness: Valid routes will be broadcasted

$$\Diamond\Diamond A \rightarrow \Diamond A$$

Topology and policies are fixed, and for liveness checks, the path needs to be provided.

Proof? How?

Proof? How?

SMT

Proof? How?

SMT

For all edge $A \rightarrow B$, generates:

$$\left(\bigwedge \text{inv}(A) \cup \text{inv}(A \rightarrow B) \cup \text{inv}(B) \right) \wedge r' = \text{Import}(A \rightarrow B, r) \\ \rightarrow \text{inv}(A \rightarrow B)[r/r']$$

Real world deployments

Hundreds of routers, tens of thousands of peering sessions.

Implemented in C#, using Zen as SMT library.

Real world deployments

Hundreds of routers, tens of thousands of peering sessions.

Implemented in C#, using Zen as SMT library.

11 configuration bugs. Each run takes at most 15 mins.

Conclusion

LIGHTYEAR is a **modular** way to verify global properties of BGP control plane.

- Guaranteed correctness for safety and liveness.
- Constraints are simple to write and understand.
- Implemented in C# which can be directly integrated into existing systems.